

PYTHON FOR HPC

- Why use python?
- Some features of the python language
- Inner workings and performance
- Multithreading and multiprocessing examples
- A few words on concurrency

Repo: <https://github.com/appeltel/python-hpc-seminar>

(repo has links, code examples, etc...)

Why Python?

WHAT IS PYTHON?

- “Scripting” Language named after Monty Python
- Created in 1989 by Guido van Rossum, who is now the Benevolent Dictator for Life
- Multi-paradigm, with object-oriented and functional features
- CPython reference implementation (current version 3.6)
- Type “import this” for language philosophy


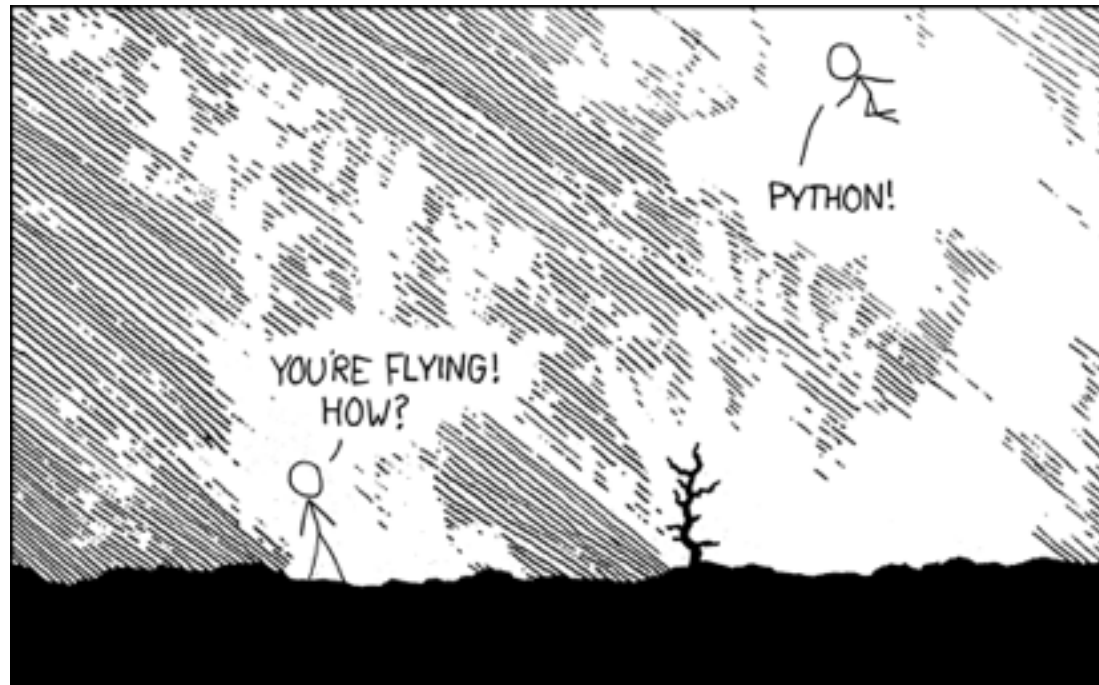


EASY TO LEARN, READ, AND USE

```
[>>> def fibonacci(n):  
[...     a, b = 0, 1  
[...     for _ in range(n):  
[...         a, b = b, a + b  
[...     return a  
[...  
[>>> fibonacci(7)  
13  
>>> █
```

- Built-in generic dynamic array (list) and mapping (dict) types
- Memory managed
- REPL, Notebooks, features for rapid prototyping
- BATTERIES INCLUDED

BATTERIES INCLUDED



- Full-featured standard library
- Extensive third-party packages
- PyPI service for distributing third party packages
- pip - tool for installing and managing packages
- virtualenv - tool for creating sandboxed environments for installing package sets



- Pre-packaged distribution of python and many scientific libraries and tools
- Tools for managing environments, installing yet more packages



django girls

PyOhio

About Python

- REPL - Read - Eval - Print - Loop
- Some basic things to explore on the REPL:
 - Names, assignment
 - Functions, modules, import system, namespaces
 - Dynamic typing, memory management
 - EAFP style - “We’re all responsible users here”
 - lists and dicts

MORE ON LISTS AND DICTS

- Built in list type “list” - (dynamic array)
- Built in mapping type “dict” - (hash table)
- Dynamic typing allows for easy representation of semi-structured data

```
>>> record = {  
...     "name": "Eric Appelt",  
...     "email": "eric.appelt@gmail.com",  
...     "uid": 111694,  
...     "dogs": ["Abby", "Rocky"]  
... }  
>>> record['dogs'][1]  
'Rocky'  
>>>
```

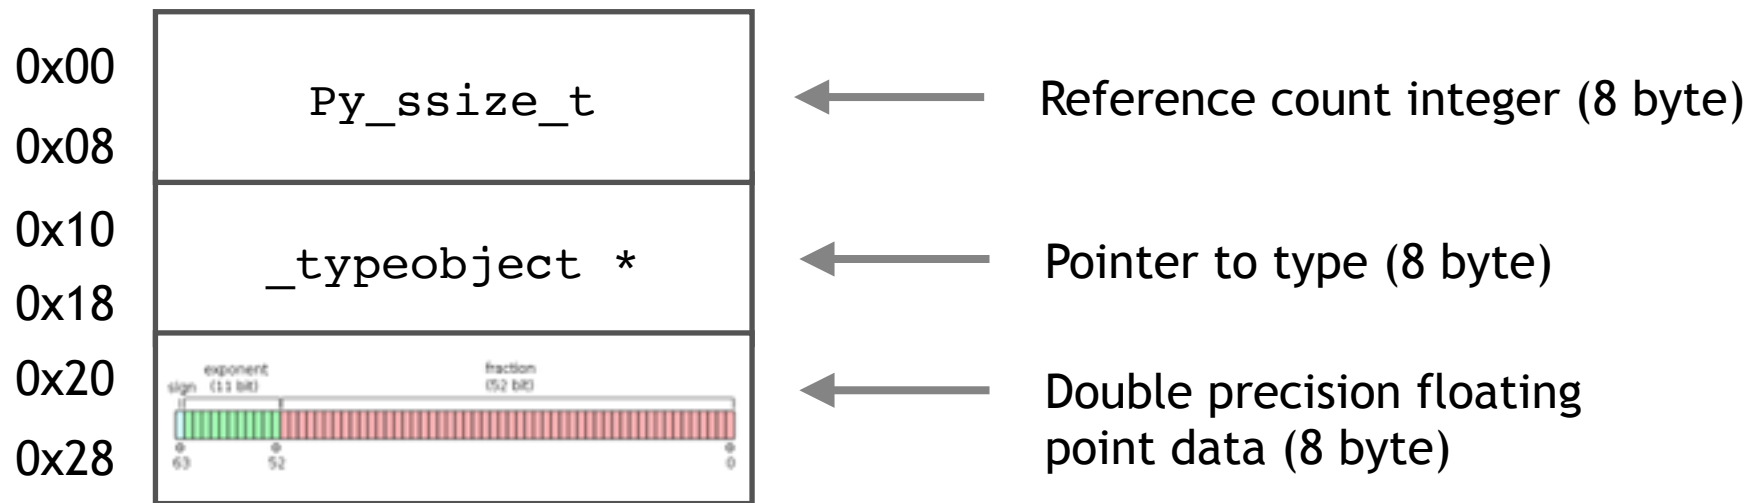
How Python Works

WARNING: This section covers the current version of the reference implementation of python called CPython. There are other implementations of the python language that do not work this way and are more performant. However, people often refer to the reference implementation when they say “python”.

- Each python object has a reference count
- Binding to a name, placement in a list position, increments the reference count
- When the reference count is zero, the object is destroyed and memory freed
- A garbage collector runs periodically to take care of circular references.

PYTHON OBJECTS IN MEMORY


Python dynamic typing and memory management by reference counting incurs memory overhead (64-bit architecture):



A python “float” object.

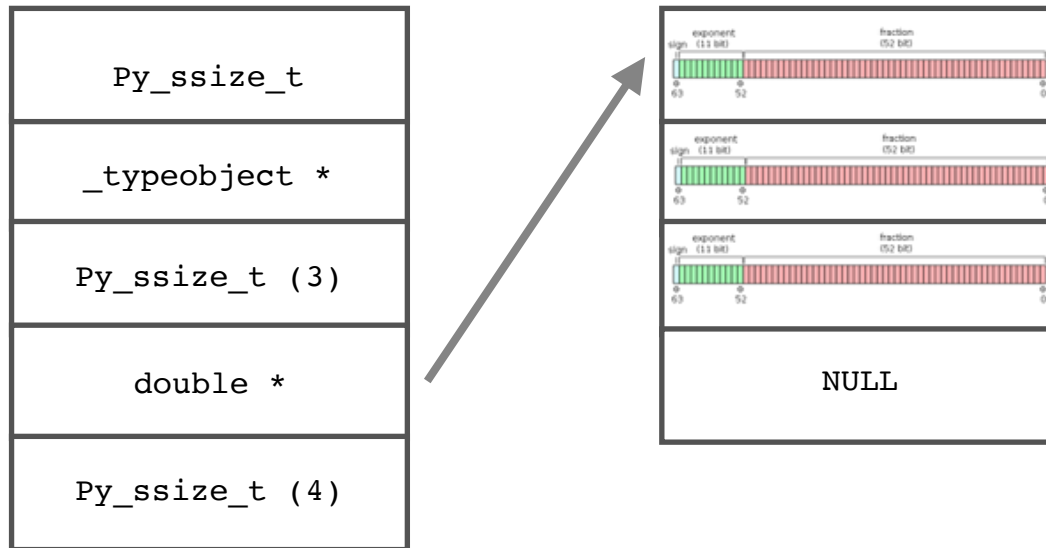
A PYTHON LIST OF FLOATS

A python list (dynamic array) of “floats”:



SOMETHING BETTER FOR HPC

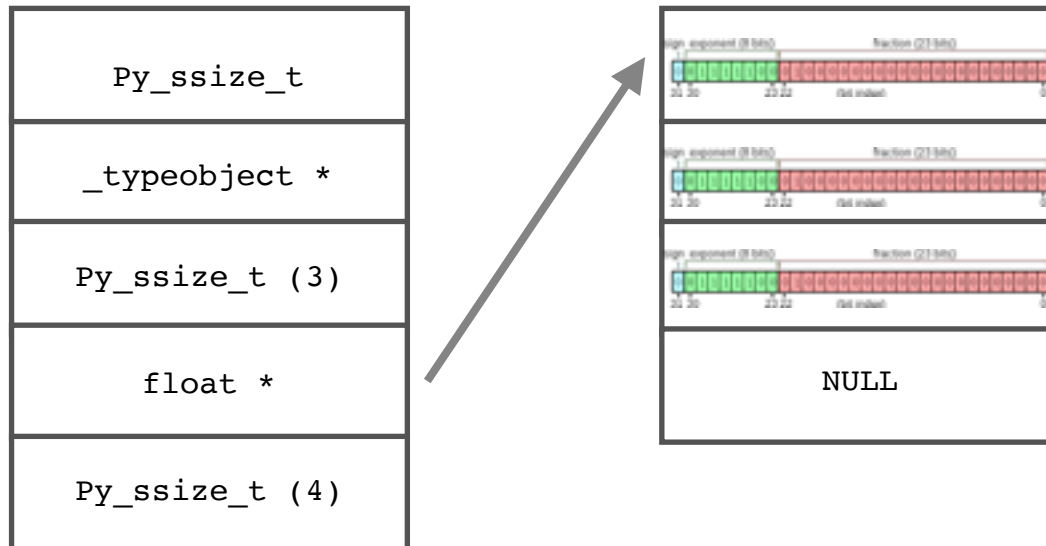
Something like this would be better for performance



8 bytes per double, no dereferencing individual pointers, checking types for each one.

SOMETHING BETTER FOR HPC

Or even better if you don't need double precision!



4 bytes per float, no
dereferencing individual
pointers, checking types
for each one.

- Python source is compiled into a platform-independent bytecode
- (CPython) A simple virtual stack machine, written in C, executes the bytecode
- (CPython) bytecode is NOT just-in-time compiled to native processor instructions
- Other implementations of python, such as pypy, do improve performance by compiling to native processor instructions.

EXAMPLE FUNCTION

```
>>> import dis
>>> def foo(x, y):
...     return x * y + 17
...
>>> foo.__code__.co_varnames
('x', 'y')
>>> foo.__code__.co_consts
(None, 17)
>>> [hex(b) for b in foo.__code__.co_code]
['0x7c', '0x0', '0x7c', '0x1', '0x14', '0x0', '0x64', '0x1', '0x17', '0x0', '0x53', '0x0']
>>> dis.dis(foo)
 2           0 LOAD_FAST           0 (x)
           2 LOAD_FAST           1 (y)
           4 BINARY_MULTIPLY
           6 LOAD_CONST           1 (17)
           8 BINARY_ADD
          10 RETURN_VALUE

>>> █
```

EXAMPLE FUNCTION INVOCATION

CALLED: `foo(3, 5)`

function code

0	LOAD_FAST	0 (x)
2	LOAD_FAST	1 (y)
4	BINARY_MULTIPLY	
6	LOAD_CONST	1 (17)
8	BINARY_ADD	
10	RETURN_VALUE	


co_varnames

PyObject * (3)	PyObject * (5)
----------------	----------------

co_consts

PyObject * (None)	PyObject * (17)
-------------------	-----------------


object stack



EXAMPLE FUNCTION INVOCATION

CALLLED: foo(3, 5)

function code



0	LOAD_FAST	0 (x)
2	LOAD_FAST	1 (y)
4	BINARY_MULTIPLY	
6	LOAD_CONST	1 (17)
8	BINARY_ADD	
10	RETURN_VALUE	

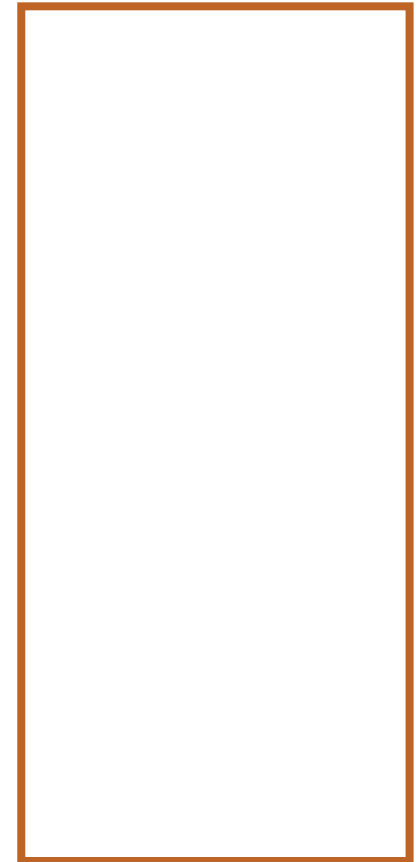
co_varnames

PyObject * (3)	PyObject * (5)
----------------	----------------

co_consts

PyObject * (None)	PyObject * (17)
-------------------	-----------------


object stack



EXAMPLE FUNCTION INVOCATION

CALLLED: foo(3, 5)

function code



0	LOAD_FAST	0 (x)
2	LOAD_FAST	1 (y)
4	BINARY_MULTIPLY	
6	LOAD_CONST	1 (17)
8	BINARY_ADD	
10	RETURN_VALUE	

co_varnames

PyObject * (3)	PyObject * (5)
----------------	----------------

co_consts

PyObject * (None)	PyObject * (17)
-------------------	-----------------

PUSH!

object stack




PyObject * (3)

EXAMPLE FUNCTION INVOCATION

CALLLED: foo(3, 5)

function code



0	LOAD_FAST	0 (x)
2	LOAD_FAST	1 (y)
4	BINARY_MULTIPLY	
6	LOAD_CONST	1 (17)
8	BINARY_ADD	
10	RETURN_VALUE	

co_varnames

PyObject * (3)	PyObject * (5)
----------------	----------------

co_consts

PyObject * (None)	PyObject * (17)
-------------------	-----------------


object stack

PyObject * (3)

EXAMPLE FUNCTION INVOCATION

CALLLED: foo(3, 5)

function code



0	LOAD_FAST	0 (x)
2	LOAD_FAST	1 (y)
4	BINARY_MULTIPLY	
6	LOAD_CONST	1 (17)
8	BINARY_ADD	
10	RETURN_VALUE	

co_varnames

PyObject * (3)	PyObject * (5)
----------------	----------------

co_consts

PyObject * (None)	PyObject * (17)
-------------------	-----------------

PUSH!

object stack




PyObject * (5)
PyObject * (3)

EXAMPLE FUNCTION INVOCATION

CALLLED: foo(3, 5)

function code



0	LOAD_FAST	0 (x)
2	LOAD_FAST	1 (y)
4	BINARY_MULTIPLY	
6	LOAD_CONST	1 (17)
8	BINARY_ADD	
10	RETURN_VALUE	

co_varnames

PyObject * (3)	PyObject * (5)
----------------	----------------

co_consts

PyObject * (None)	PyObject * (17)
-------------------	-----------------

object stack

PyObject * (5)


PyObject * (3)

EXAMPLE FUNCTION INVOCATION

CALLER: foo(3, 5)

function code

0	LOAD_FAST	0 (x)
2	LOAD_FAST	1 (y)
4	BINARY_MULTIPLY	
6	LOAD_CONST	1 (17)
8	BINARY_ADD	
10	RETURN_VALUE	



co_varnames

PyObject * (3)	PyObject * (5)
----------------	----------------

co_consts

PyObject * (None)	PyObject * (17)
-------------------	-----------------

object stack

PyObject * (5)

POP!




PyObject * (3)

EXAMPLE FUNCTION INVOCATION

CALLLED: foo(3, 5)

function code



0	LOAD_FAST	0 (x)
2	LOAD_FAST	1 (y)
4	BINARY_MULTIPLY	
6	LOAD_CONST	1 (17)
8	BINARY_ADD	
10	RETURN_VALUE	

co_varnames

PyObject * (3)	PyObject * (5)
----------------	----------------

co_consts

PyObject * (None)	PyObject * (17)
-------------------	-----------------

object stack

PyObject * (5)

PyObject * (3)

POP !




EXAMPLE FUNCTION INVOCATION

CALLLED: `foo(3, 5)`

function code

0	LOAD_FAST	0 (x)
2	LOAD_FAST	1 (y)
4	BINARY_MULTIPLY	
6	LOAD_CONST	1 (17)
8	BINARY_ADD	
10	RETURN_VALUE	



co_varnames

PyObject * (3)	PyObject * (5)
----------------	----------------

co_consts

PyObject * (None)	PyObject * (17)
-------------------	-----------------

`__mult__`

PyObject * (5)

PyObject * (3)



PyObject * (15)

object stack




EXAMPLE FUNCTION INVOCATION

CALLLED: foo(3, 5)

function code

0	LOAD_FAST	0 (x)
2	LOAD_FAST	1 (y)
4	BINARY_MULTIPLY	
6	LOAD_CONST	1 (17)
8	BINARY_ADD	
10	RETURN_VALUE	



co_varnames

PyObject * (3)	PyObject * (5)
----------------	----------------

co_consts

PyObject * (None)	PyObject * (17)
-------------------	-----------------

object stack

PyObject * (15)

PyObject * (15)




PUSH!

EXAMPLE FUNCTION INVOCATION

CALLER: foo(3, 5)

function code



0	LOAD_FAST	0 (x)
2	LOAD_FAST	1 (y)
4	BINARY_MULTIPLY	
6	LOAD_CONST	1 (17)
8	BINARY_ADD	
10	RETURN_VALUE	

co_varnames

PyObject * (3)	PyObject * (5)
----------------	----------------

co_consts

PyObject * (None)	PyObject * (17)
-------------------	-----------------


object stack

PyObject * (15)

EXAMPLE FUNCTION INVOCATION

CALLLED: `foo(3, 5)`

function code



0	LOAD_FAST	0 (x)
2	LOAD_FAST	1 (y)
4	BINARY_MULTIPLY	
6	LOAD_CONST	1 (17)
8	BINARY_ADD	
10	RETURN_VALUE	


co_varnames

PyObject * (3)	PyObject * (5)
----------------	----------------

co_consts

PyObject * (None)	PyObject * (17)
-------------------	-----------------

object stack



PUSH!


PyObject * (17)
PyObject * (15)

EXAMPLE FUNCTION INVOCATION

CALLLED: foo(3, 5)

function code

0	LOAD_FAST	0 (x)
2	LOAD_FAST	1 (y)
4	BINARY_MULTIPLY	
6	LOAD_CONST	1 (17)
8	BINARY_ADD	
10	RETURN_VALUE	



co_varnames

PyObject * (3)	PyObject * (5)
----------------	----------------

co_consts

PyObject * (None)	PyObject * (17)
-------------------	-----------------

object stack


PyObject * (17)
PyObject * (15)

EXAMPLE FUNCTION INVOCATION

CALLLED: `foo(3, 5)`

function code

0	LOAD_FAST	0 (x)
2	LOAD_FAST	1 (y)
4	BINARY_MULTIPLY	
6	LOAD_CONST	1 (17)
8	BINARY_ADD	
10	RETURN_VALUE	



co_varnames

PyObject * (3)	PyObject * (5)
----------------	----------------

co_consts

PyObject * (None)	PyObject * (17)
-------------------	-----------------

object stack

PyObject * (17)

POP!

PyObject * (15)




EXAMPLE FUNCTION INVOCATION

CALLLED: foo(3, 5)

function code

0	LOAD_FAST	0 (x)
2	LOAD_FAST	1 (y)
4	BINARY_MULTIPLY	
6	LOAD_CONST	1 (17)
8	BINARY_ADD	
10	RETURN_VALUE	



co_varnames

PyObject * (3)	PyObject * (5)
----------------	----------------

co_consts

PyObject * (None)	PyObject * (17)
-------------------	-----------------

object stack

PyObject * (17)

PyObject * (15)

POP!




EXAMPLE FUNCTION INVOCATION

CALLLED: foo(3, 5)

function code

0	LOAD_FAST	0 (x)
2	LOAD_FAST	1 (y)
4	BINARY_MULTIPLY	
6	LOAD_CONST	1 (17)
8	BINARY_ADD	
10	RETURN_VALUE	



co_varnames

PyObject * (3)	PyObject * (5)
----------------	----------------

co_consts


PyObject * (None)	PyObject * (17)
-------------------	-----------------

object stack

__add__

PyObject * (17)
PyObject * (15)

PyObject * (32)




EXAMPLE FUNCTION INVOCATION

CALLLED: foo(3, 5)

function code

0	LOAD_FAST	0 (x)
2	LOAD_FAST	1 (y)
4	BINARY_MULTIPLY	
6	LOAD_CONST	1 (17)
8	BINARY_ADD	
10	RETURN_VALUE	



co_varnames

PyObject * (3)	PyObject * (5)
----------------	----------------

co_consts

PyObject * (None)	PyObject * (17)
-------------------	-----------------

object stack

PyObject * (32)

PyObject * (32)


PUSH!

EXAMPLE FUNCTION INVOCATION

CALLLED: foo(3, 5)

function code

0	LOAD_FAST	0 (x)
2	LOAD_FAST	1 (y)
4	BINARY_MULTIPLY	
6	LOAD_CONST	1 (17)
8	BINARY_ADD	
10	RETURN_VALUE	



co_varnames

PyObject * (3)	PyObject * (5)
----------------	----------------

co_consts

PyObject * (None)	PyObject * (17)
-------------------	-----------------

object stack


PyObject * (32)

EXAMPLE FUNCTION INVOCATION

CALLLED: foo(3, 5)

function code

0	LOAD_FAST	0 (x)
2	LOAD_FAST	1 (y)
4	BINARY_MULTIPLY	
6	LOAD_CONST	1 (17)
8	BINARY_ADD	
10	RETURN_VALUE	



co_varnames

PyObject * (3)	PyObject * (5)
----------------	----------------

co_consts

PyObject * (None)	PyObject * (17)
-------------------	-----------------

object stack

PyObject * (32)

POP!

PyObject * (32)




EXAMPLE FUNCTION INVOCATION

CALLLED: `foo(3, 5)`

function code

0	LOAD_FAST	0 (x)
2	LOAD_FAST	1 (y)
4	BINARY_MULTIPLY	
6	LOAD_CONST	1 (17)
8	BINARY_ADD	
10	RETURN_VALUE	



co_varnames

PyObject * (3)	PyObject * (5)
----------------	----------------

co_consts

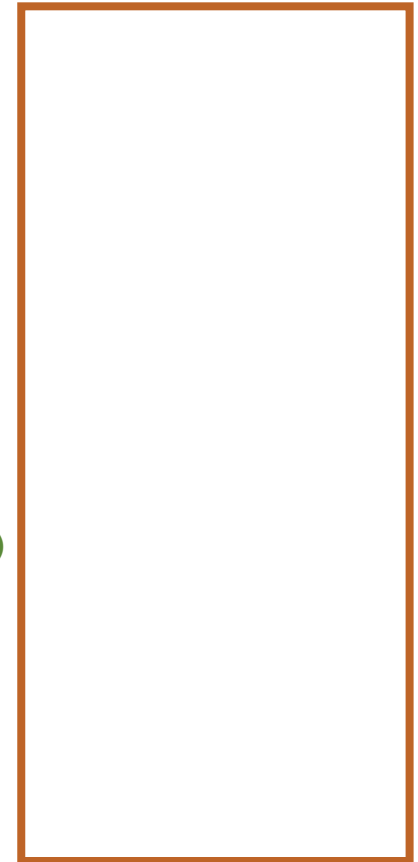
PyObject * (None)	PyObject * (17)
-------------------	-----------------

(return to caller)

PyObject * (32)



object stack



Parallelism in Python

- Take some LU-decompositon and inversion example code written in pure python
- Run it serially over a set of random matrices
- Try improving performance by multithreading
- Try improving performance by multiprocessing

Example code in the repo if you (or I) fall behind!

About Concurrency

- Parallelism: Do multiple things at the same time
- Concurrency: Deal with multiple things going on at the same time
- Cooking example

Python supports multiple paradigms for achieving concurrent programming including os-level threads and also native coroutines

PYTHON COROUTINE EXAMPLE

```
import asyncio
```

```
import aiohttp
```

```
BASE_URL = 'https://ericappelt.com/animals/'
```

```
async def speak(animal, session):
```

```
    async with session.get('{}{}'.format(BASE_URL, animal)) as response:
        response.raise_for_status()
        sound = await response.text()
```

```
    return 'The {} says "{}".'.format(animal, sound)
```

Day Two

- Packaging, pip, and anaconda
- Numpy
- Revisiting matrix inversions with numpy
- A quick tour of scientific packages
- Jupyter notebooks and literate programming
- Making pretty and interactive plots

Repo: <https://github.com/appeltel/python-hpc-seminar>

(repo has links, code examples, etc...)

External Libraries

- pip: download and additional python libraries
- pypi/warehouse: central repository for community libraries
- virtualenv: tool for managing lightweight python “environments” of installed libraries
- system python: Python installed as part of your operating system - you don't want to mess with this

Quick Demo (... or read below if something goes wrong ...)

```
eric@laptop:tmp$ python -m virtualenv venv
Using base prefix '/Users/eric/.pyenv/versions/3.6.2'
New python executable in /Users/eric/tmp/venv/bin/python
Installing setuptools, pip, wheel...done.
eric@laptop:tmp$ source venv/bin/activate
(venv) eric@laptop:tmp$ pip install numpy
Collecting numpy
  Downloading numpy-1.14.2-cp36-cp36m-macosx_10_6_intel.macosx_10_9_intel.ma
ntel.macosx_10_10_x86_64.whl (4.7MB)
    100% |██████████████████████████████████████████████████████████████████| 4.7MB 193kB/s
Installing collected packages: numpy
Successfully installed numpy-1.14.2
(venv) eric@laptop:tmp$ python -c "import numpy as np; print(np.arange(5))"
[0 1 2 3 4]
(venv) eric@laptop:tmp$ deactivate
eric@laptop:tmp$ rm -r venv
eric@laptop:tmp$ █
```



- Pre-packaged distribution of python and many scientific libraries and tools
- Tools for managing environments, installing yet more packages
- Available on ACCRE cluster via LMOD

Intro to Numpy



NumPy is the fundamental package for scientific computing with Python. It contains among other things:

- a powerful N-dimensional array object
- sophisticated (broadcasting) functions
- tools for integrating C/C++ and Fortran code
- useful linear algebra, Fourier transform, and random number capabilities

Besides its obvious scientific uses, NumPy can also be used as an efficient multi-dimensional container of generic data. Arbitrary data-types can be defined. This allows NumPy to seamlessly and speedily integrate with a wide variety of databases.

Quickstart:

<https://docs.scipy.org/doc/numpy-dev/user/quickstart.html>

- Some experimentation on the REPL
- Revisiting the matrix inversion code
- Look at multithreading performance




Tour of Libraries



- The SciPy library provides user friendly numerical integration and optimization functions
- Built on NumPy
- [SciPy.org](https://www.scipy.org) is the organization maintaining numpy, scipy, and a set of related tools
- These tools are often referred to as the SciPy stack



matplotlib

Matplotlib is a Python 2D plotting library which produces publication quality figures in a variety of hardcopy formats and interactive environments across platforms. Matplotlib can be used in Python scripts, the Python and **IPython** shells, the **Jupyter** notebook, web application servers, and four graphical user interface toolkits.



pandas


$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$



pandas is a **Python** package providing fast, flexible, and expressive data structures designed to make working with "relational" or "labeled" data both easy and intuitive. It aims to be the fundamental high-level building block for doing practical, **real world** data analysis in Python. Additionally, it has the broader goal of becoming **the most powerful and flexible open source data analysis / manipulation tool available in any language**. It is already well on its way toward this goal.

pandas is well suited for many different kinds of data:

- Tabular data with heterogeneously-typed columns, as in an SQL table or Excel spreadsheet
- Ordered and unordered (not necessarily fixed-frequency) time series data.
- Arbitrary matrix data (homogeneously typed or heterogeneous) with row and column labels
- Any other form of observational / statistical data sets. The data actually need not be labeled at all to be placed into a pandas data structure

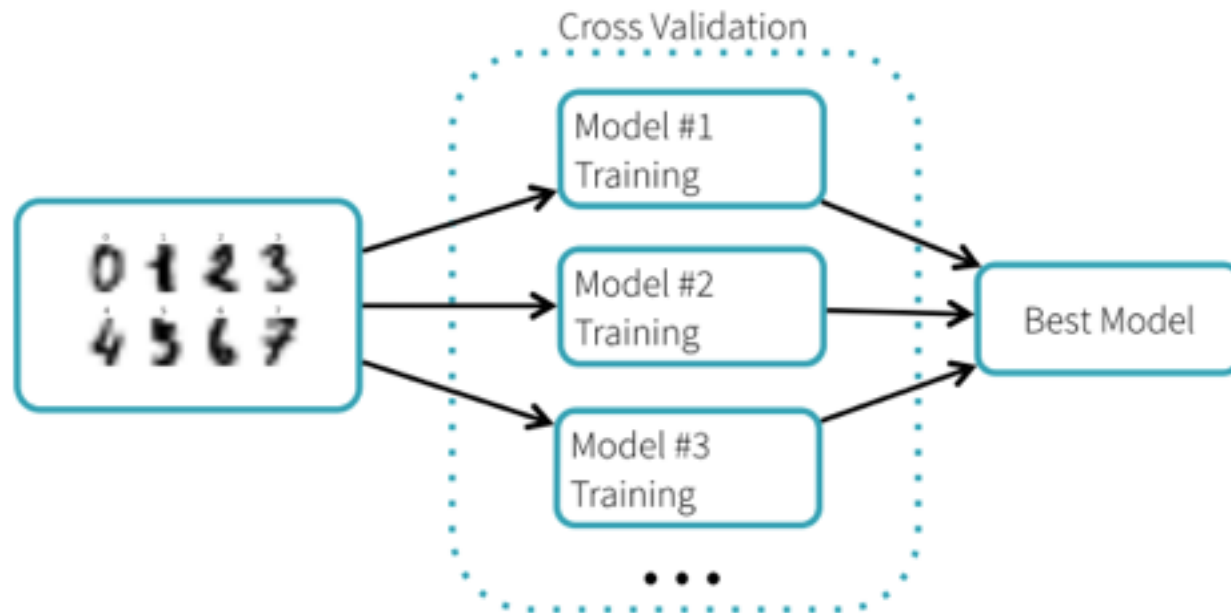


scikit-learn

Machine Learning in Python

- Simple and efficient tools for data mining and data analysis
- Accessible to everybody, and reusable in various contexts
- Built on NumPy, SciPy, and matplotlib
- Open source, commercially usable - BSD license

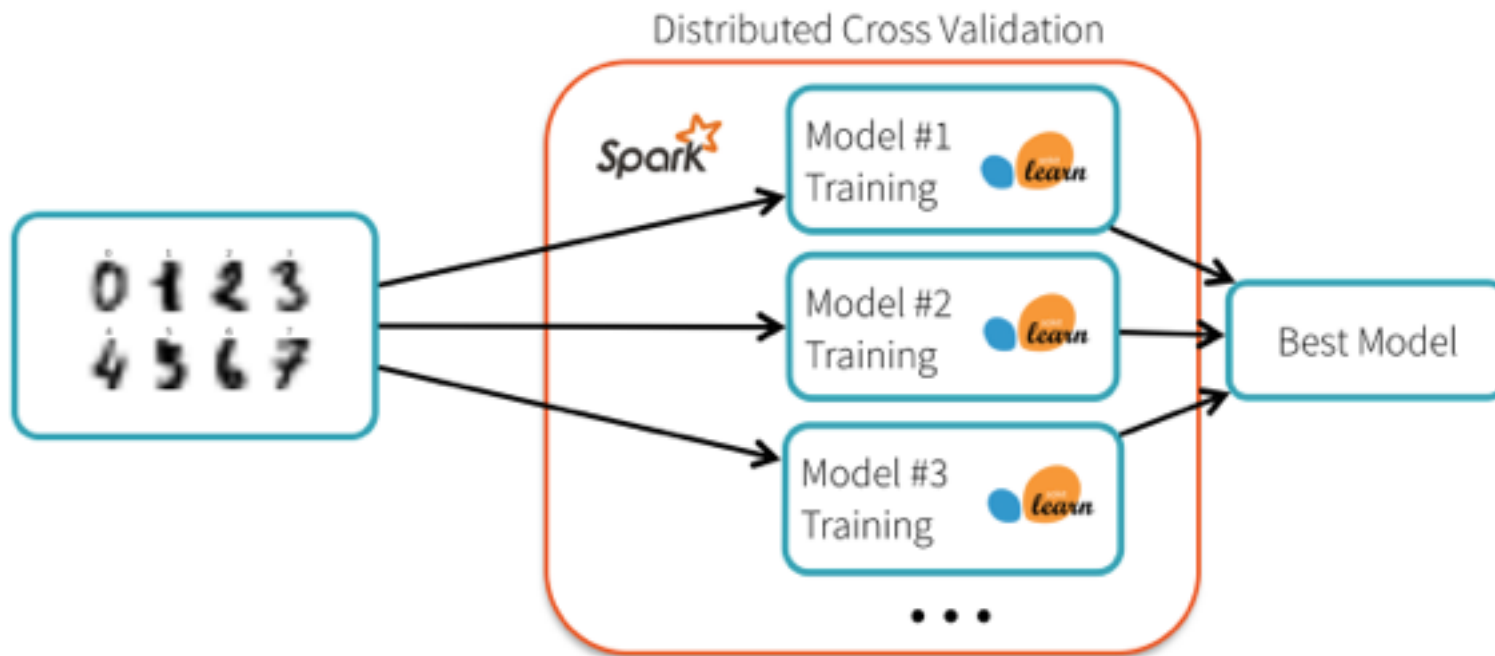
Training a random forest classifier to recognize handwritten digits:



Good candidate for parallelization and distribution over multiple cluster nodes, see:

<https://databricks.com/blog/2016/02/08/auto-scaling-scikit-learn-with-apache-spark.html>

Easy to do on a spark cluster:



- Can it also be done within a slurm job that sets up a joblib cluster for the training?
- Can we use our big data cluster for this?
- How is ACCRE performance?

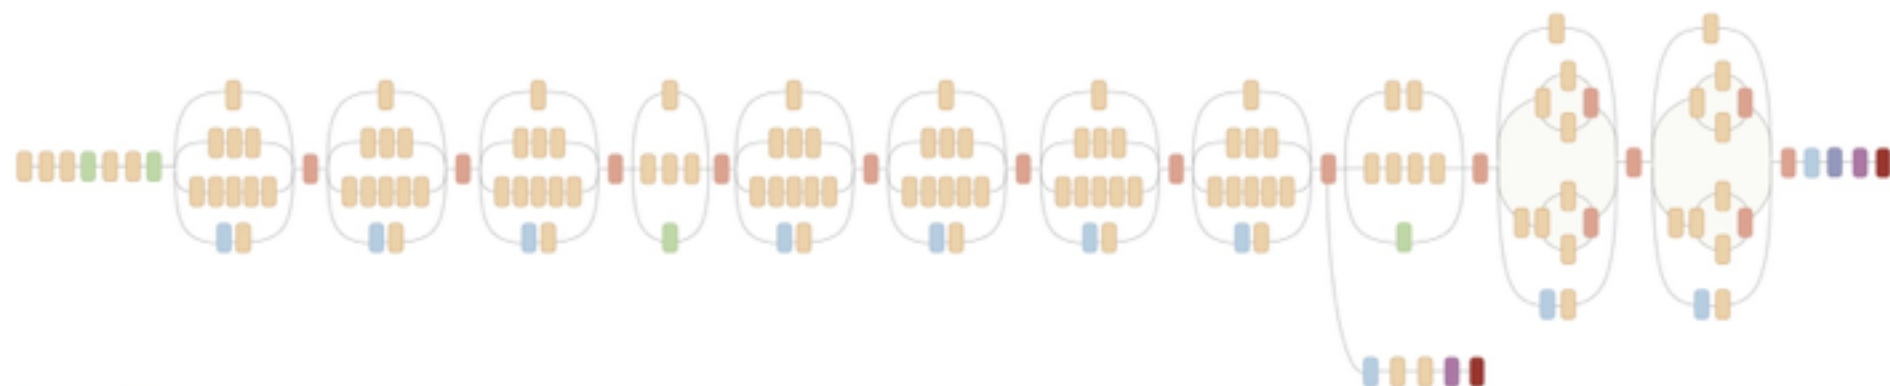


- Numerical computation with data flow graphs
- Edges represent mathematical operations
- Vertices represent tensors (arrays, matrices)
- Common API for deployment to CPU/GPU
- Deep Learning applications

Fun demo of deep learning: playground.tensorflow.org

TENSORFLOW - IMAGE RECOGNITION

Inception-v3 model:



- Convolution
- AvgPool
- MaxPool
- Concat
- Dropout
- Fully connected
- Softmax



Jupyter Notebooks



- Interactive web application
- Embed live code, narrative, and visualizations into notebook files
- Can be run from the ACCRE cluster
- Not just python!



- Do some demos
- You can run these directly out of the repo:
- <https://github.com/appeltel/python-hpc-seminar>